

OSQP Settings Sequence Data Frame

Data frame containing a sequence of OSQP settings for `tsbalancing()` specified with argument `osqp_settings_df`. The package includes two predefined OSQP settings sequence data frames that are presented below.

Data frame `default_osqp_sequence`

Fast and effective sequence of OSQP settings that should be suitable for accurately solving most time series balancing problems. It is the default value of `tsbalancing()` argument `osqp_settings_df`.

```
#>   max_iter   sigma eps_abs eps_rel eps_prim_inf eps_dual_inf polish scaling
#> 1    4,000 1.00e-09 1.00e-06 1.00e-06    1.00e-07    1.00e-07  TRUE      0
#> 2   10,000 1.00e-15 1.00e-12 1.00e-12    1.00e-13    1.00e-13  TRUE      0
#> 3   10,000 1.00e-15 1.00e-12 1.00e-12    1.00e-13    1.00e-13  TRUE      0
#> 4   10,000 2.22e-16 2.22e-16 2.22e-16    2.22e-16    2.22e-16  TRUE      0
#>   prior_scaling require_polished
#> 1           TRUE              TRUE
#> 2           TRUE              TRUE
#> 3          FALSE              FALSE
#> 4           TRUE              FALSE
```

Data frame `alternate_osqp_sequence`

Alternative slower sequence of OSQP settings that could help achieve more precision, when needed, especially when combined with argument `full_sequence = TRUE`.

```
#>   max_iter   sigma eps_abs eps_rel eps_prim_inf eps_dual_inf polish scaling
#> 1   10,000 1.00e-15 1.00e-12 1.00e-12    1.00e-13    1.00e-13  TRUE      0
#> 2   10,000 1.00e-15 1.00e-12 1.00e-12    1.00e-13    1.00e-13  TRUE     10
#> 3   10,000 1.00e-12 1.00e-09 1.00e-09    1.00e-10    1.00e-10  TRUE      0
#> 4   10,000 1.00e-12 1.00e-09 1.00e-09    1.00e-10    1.00e-10  TRUE     10
#> 5   10,000 1.00e-09 1.00e-06 1.00e-06    1.00e-07    1.00e-07  TRUE      0
#> 6   10,000 1.00e-09 1.00e-06 1.00e-06    1.00e-07    1.00e-07  TRUE     10
#> 7   10,000 1.00e-06 1.00e-03 1.00e-03    1.00e-04    1.00e-04  TRUE      0
#> 8   10,000 1.00e-06 1.00e-03 1.00e-03    1.00e-04    1.00e-04  TRUE     10
#> 9   10,000 2.22e-16 2.22e-16 2.22e-16    2.22e-16    2.22e-16  TRUE      0
#> 10  10,000 2.22e-16 2.22e-16 2.22e-16    2.22e-16    2.22e-16  TRUE     10
#> 11  10,000 1.00e-15 1.00e-12 1.00e-12    1.00e-13    1.00e-13  TRUE      0
#> 12  10,000 1.00e-12 1.00e-09 1.00e-09    1.00e-10    1.00e-10  TRUE      0
#> 13  10,000 1.00e-09 1.00e-06 1.00e-06    1.00e-07    1.00e-07  TRUE      0
#> 14  10,000 1.00e-06 1.00e-03 1.00e-03    1.00e-04    1.00e-04  TRUE      0
#> 15  10,000 2.22e-16 2.22e-16 2.22e-16    2.22e-16    2.22e-16  TRUE      0
#>   prior_scaling require_polished
#> 1          FALSE              TRUE
#> 2          FALSE              TRUE
#> 3          FALSE              TRUE
#> 4          FALSE              TRUE
#> 5          FALSE              TRUE
#> 6          FALSE              TRUE
```

#> 7	FALSE	TRUE
#> 8	FALSE	TRUE
#> 9	FALSE	TRUE
#> 10	FALSE	TRUE
#> 11	TRUE	TRUE
#> 12	TRUE	TRUE
#> 13	TRUE	TRUE
#> 14	TRUE	TRUE
#> 15	TRUE	TRUE

Details

With the exception of `prior_scaling` and `require_polished`, all columns of the data frame must correspond to a OSQP setting. Default OSQP values are used for any setting not specified in this data frame. Visit https://osqp.org/docs/interfaces/solver_settings.html for all available OSQP settings. Note that the OSQP `verbose` setting is actually controlled through `tsbalancing()` arguments `quiet` and `display_level` (i.e., column `verbose` in a *OSQP settings sequence data frame* would be ignored).

Each row of a *OSQP settings sequence data frame* represents one attempt at solving a balancing problem with the corresponding OSQP settings. The solving sequence stops as soon as a valid solution is obtained (a solution for which all constraint discrepancies are smaller or equal to the tolerance specified with `tsbalancing()` argument `validation_tol`) unless column `require_polished` = `TRUE`, in which case a polished solution from OSQP (`status_polish` = 1) would also be required to stop the sequence. Constraint discrepancies correspond to $\max(0, l - Ax, Ax - u)$ with constraints defined as $l \leq Ax \leq u$. In the event where a satisfactory solution cannot be obtained after having gone through the entire sequence, `tsbalancing()` returns the solution that generated the smallest total constraint discrepancies among valid solutions, if any, or among all solutions, otherwise. Note that running the entire solving sequence can be *enforced* by specifying `tsbalancing()` argument `full_sequence` = `TRUE`. Rows with column `prior_scaling` = `TRUE` have the problem data scaled prior to solving with OSQP, using the average of the free (nonbinding) problem values as the scaling factor.

In addition to specifying a custom-made *OSQP settings sequence data frame* with argument `osqp_settings_df`, one can also specify `osqp_settings_df` = `NULL` which would result in a single solving attempt with default OSQP values for all settings along with `prior_scaling` = `FALSE` and `require_polished` = `FALSE`. Note that it is recommended, however, to first try data frames `default_osqp_sequence` and `alternate_osqp_sequence`, along with `full_sequence` = `TRUE` if necessary, before considering other alternatives.

Recommended Approach

Start with the default `tsbalancing()` solving sequence (`osqp_settings_df` = `default_osqp_sequence` and `full_sequence` = `FALSE`). Then, if more precision is needed, try with:

1. `full_sequence` = `TRUE`
2. `osqp_settings_df` = `alternate_osqp_sequence`
3. `osqp_settings_df` = `alternate_osqp_sequence` and `full_sequence` = `TRUE`

In practice, specifying `full_sequence` = `TRUE` should be enough when more precision is needed (at the expense of execution time, obviously). Only in rare occasions should you need to use the `alternate_osqp_sequence` data frame, which will often be even more costly in terms of execution time especially when combined with `full_sequence` = `TRUE`.

Guiding Principles

The following is a summary of the lessons learned while developing `tsbalancing()` and experimenting with the OSQP solver. They are the guiding principles that lead to both *OSQP settings sequence data*

frames presented earlier. Note these observations apply to **time series balancing problems** as solved by `tsbalancing()` and may not directly apply to other types of *quadratic problems*.

- Data preconditioning options available in OSQP (with the `scaling` setting) are not sufficient for some (badly scaled) problems. External (prior) data scaling (`prior_scaling = TRUE`) is sometimes necessary for OSQP to converge at a decent pace and generate *precise enough* solutions in a reasonable number of iterations.
- Prior data scaling often reduces execution time (the required number of iterations to achieve the specified precision) and greatly increases the likelihood of polished solutions.
- Polished solutions are always very precise, even when prior data scaling is performed (i.e., the solution in the original scale will usually still be *precise enough*).
- While polished solutions from prior-scaled data are usually more precise than unpolished solutions from non prior-scaled data, the most precise solutions correspond to polished solutions from non prior-scaled data.
- Smaller `sigma` and tolerance (`eps_*`) settings result in more precise solutions but take longer to run (require more iterations).
- Enough precision is usually obtained after 10,000 iterations with small values for the `sigma` and tolerance (`eps_*`) settings.
- The default OSQP values for `alpha` and the various settings associated to ρ (`*rho*`) are sufficient (they work well). Reducing the `sigma` and tolerance (`eps_*`) settings and performing prior data scaling is sufficient to obtain precise solutions in a reasonable number of iterations.
- Keeping the same scale between the `sigma` and tolerance (`eps_*`) settings as the default OSQP values works well and is used in both *OSQP settings sequence data frames*, i.e.:
 - `eps_abs = eps_rel = 1,000 * sigma`
 - `eps_prim_inf = eps_dual_inf = 100 * sigma`
 - (and consequently) `eps_abs = eps_rel = 10 * eps_prim_inf = 10 * eps_dual_inf`
- The *machine epsilon* (`.Machine$double.eps`) for the `sigma` and tolerance (`eps_*`) settings, which basically forces the maximum number of iterations, is used as a *last resort* in both *OSQP settings sequence data frames*.

Summary - default sequence (data frame `default_osqp_sequence`)

- Geared towards achieving both fast and precise solutions.
- First try to get (fast) polished solutions on prior-scaled data before attempting solving the original (non-scaled) problem data.
- Make a final attempt on prior-scaled data with the *machine epsilon* for the `sigma` and tolerance (`eps_*`) settings.

Summary - alternative sequence (data frame `alternate_osqp_sequence`)

- Geared towards achieving precise solutions at the expense of execution time.
- Somewhat similar to a *brute force* or *try all* approach (especially when combined with `full_sequence = TRUE`).
- Small `sigma` and tolerance (`eps_*`) settings that gradually increase, with the *machine epsilon* as a last attempt.
- Polished solutions are required for every step of the sequence (the best unpolished solution is returned if no polished solution could be obtained at the end of the entire sequence).
- Maximum of 10,000 iterations for every step of the sequence
- First try to get polished solutions from non prior-scaled data (the most precise solutions), then try with prior-scaled data.